AI Cheat-sheet 2025 July Edition

No code automation tools

No-code automation tools let you connect apps like Gmail, Google Sheets, WhatsApp, or Slack and automate repetitive tasks without writing any code. It's like drag and drop nodes, giving a visual representation of your workflow. For example, sending a message when a form is filled or auto updating a spreadsheet when a payment is made. They save time, reduce manual work, and let anyone, even non-developers, build smart, customized automations for their business or personal use.

Platform Name	Best Use Cases	Ease of Use	Developer Flexibility	Pricing Model	Hosting Options
n8n	Agentic workflows, API automations, backend bots, self-hosted agents	Moderate (Visual UI + dev-friendly, but some setup needed)	High (Custom JS/HTTP nodes, full data access, module support)	Free (OSS) + cloud plans (starts ~$20/month)	Self-hosted & Cloud
Make	Visually rich no-code automation, eComm ops, marketing tools	High (Drag-and-drop UI, templates, less technical)	Moderate (Basic HTTP module, but limited code execution)	Free tier + usage-based paid plans	Cloud-only
Pipedream	API integrations, event-driven dev workflows, SaaS middleware	Moderate (CLI-like UI, best for developers)	Very High (Run JS, Python, custom packages)	Free tier + pay-per-invocation (dev-friendly)	Cloud-only
Zapier	Simple workflows, no-code business automation, CRM triggers	Very High (No-code first, easy templates)	Low (limited code blocks, constrained HTTP support)	Tiered plans; higher tiers expensive for scale	Cloud-only

Text Generation Models (Chat, Coding, Summarization etc)

Text generation models are a type of AI that can read, understand, and write human-like text. You give them a prompt, like a question, a sentence, or even a long document and they generate a response that sounds natural and intelligent. They're trained on huge amounts of data from books, websites, and conversations, so they've picked up how we communicate. These models can help with everything from writing emails and summarizing reports to coding, translating languages, or chatting like a virtual assistant.

Model	Release & Developer	Modalities & Context Window	Benchmark Performance	Pricing & Access	Best Use Case
Grok 4	July 2025 – xAI (Elon Musk)	Text, code, web browsing; limited vision support coming; context window: not disclosed	State-of-the-art in GPQA (100%), AIME (100%), and grad-level physics (87%)	$30/month (X Premium+), $300/month (Heavy tier); API: ~$3 input / $15 output per 1M tokens	Real-time data tasks (e.g., financial markets, breaking news), live-coding, deep math reasoning

Rahul Kamath (https://www.linkedin.com/in/rahulnkamath/)

GPT-4o	May 2024 – OpenAI	Text, vision, audio (input/output); 128K token context window	MMLU: 88.7; top-tier in vision, math, multilingual tasks	ChatGPT Plus: $20/month; API: $2.50 input / $10 output per 1M tokens	Multimodal agents, creative content, audio/image apps, low-latency UX
LLaMA 4	April 2025 – Meta AI	Text + image (Scout/Maverick/Behemoth); context up to 10M tokens (Scout), >1M (Behemoth)	Maverick exceeds GPT-4o in some logic & coding; Behemoth > Claude/GPT in STEM	Open-source (free); API (e.g., Together.ai): ~$0.18 input / $0.85 output per 1M tokens	Self-hosted research apps, massive document analysis, budget-friendly scale
Claude Opus 4	May 2025 – Anthropic	Text, code, reasoning; hybrid "Opus + Haiku" runtime; context est. 200K+ (not public)	SWE-bench: 72.5% (vs GPT-4.1: 54.6%); excels in long-context coding	Free tier; Claude Pro: $17/month; API: $15 input / $75 output per 1M tokens	Sustained reasoning, long-form coding, multi-hour research use cases

Image Generation Models

Image generation models are AI systems that can turn text prompts into picture. You just describe what you want, like "a futuristic city at sunset" or "a cat dressed as a chef," and the model creates a visual version of it. Some models are great at making realistic photos, others are better for illustrations, design work, or fantasy art. Over the last year, they've gotten crazy good, not just at making things look cool, but also at handling details like lighting, perspective, or even adding text in specific fonts.

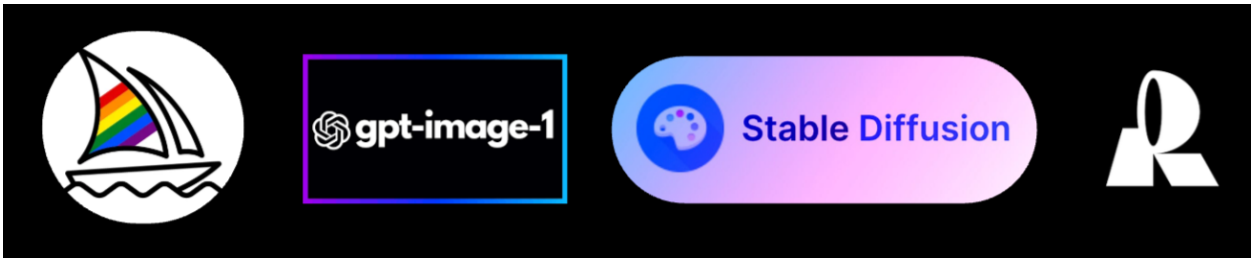

Model	Core Features	Strengths	Weaknesses	Open Source	Pricing	Best Use Cases
Midjourney V7	High-quality art, personalization, draft mode, style remixing	Most artistic & visually striking outputs; community-rich	Closed model, no API, lacks exact object/text accuracy	No	Subscription-based; Turbo credits consume faster	Concept art, illustrations, mood boards, fantasy art
GPT-Image-1	Text + image input, multimodal edits, GPT-4 integration	Excellent prompt fidelity, image+text accuracy, enterprise integrations	Pay-per-use, closed-source, limited artistic flair by default	No	Pay-per-image (API); included in ChatGPT Plus tier	Marketing, UI design, brand graphics, in-chat visual tasks
Stable Diffusion 3.5	Open-source, multiple model sizes, inpainting, community finetunes	Customizable, efficient, great for photorealism and R&D	Needs prompt tuning; no guardrails; requires good GPU for best version	Yes	Free (under open license); optional enterprise license >$1M revenue	Photorealism, science/medical visualization, anime/art styles
Recraft V3	Precise layout control, long text rendering, vector outputs, design-focused tools	Text + layout control, best for commercial design, brand consistency	Closed ecosystem, fewer integrations, learning curve for advanced features	No	Free tier with paid plans; API metered for commercial use	Posters, ads, icon sets, flyers, any image with embedded text

Video Generation Model

Video generation models are AI systems that can create entire videos from scratch using just a text prompt, a still image, or a short video clip as input. These models work by understanding what you're asking for (like "a lion running in the desert during sunset") and then generating multiple frames that visually tell that story — complete with lighting, movement, and sometimes even cinematic camera angles or sound. Some models also let you control the motion style, camera path, or emotional tone of characters.

Rahul Kamath (https://www.linkedin.com/in/rahulnkamath/)

Model Name	Provider	Best Use Cases	Model Strength	Commercial Availability	Control Features	Pricing
Veo 3	Google DeepMind	High-quality cinematic video, storytelling, ads	Long-form video generation, camera control, text-to-video	Limited beta, YouTube Shorts partners only (as of 2025)	Prompt + cinematic tools (camera, scene, composition)	Not publicly available; invite-only
Kling AI	Kuaishou (China)	Short-form content, TikTok-style videos, influencers	Realistic human motion, lip sync, emotional expressions	Closed beta, Chinese creators (planned public access)	Text + motion reference + face video input	Free (in beta); no commercial tiers disclosed
Hualia	Huya (China)	AI streamers, avatars, VTubing, real-time deepfakes	Real-time AI avatars, facial mimicry, stream integration	China-only, partnership-based	Real-time control, face/voice mimicry, emotion presets	B2B licensing; no public pricing

Voice Generation Models

Voice generation models are basically AI tools that can talk — and not in that old robotic way. These models can mimic human speech with surprisingly natural tone, pacing, and even emotion. You can type out a script, and the AI reads it out loud in a voice that sounds like a real person, sometimes even in different accents or styles. Some models can clone your voice, others let you pick from a wide variety of characters or tones — from calm narrators to energetic presenters.

Model	Type	Best Use Cases	Voice Quality	Multilingual Support	Pricing	Integration
Google Cloud TTS	Commercial	IVR, global apps, accessibility, e-learning	Very Good (Neural2/WaveNet)	50+ languages	$16/M chars (Neural), $4/M (Standard)	Easy API/SDK, Cloud-only
Microsoft Azure TTS	Commercial	Enterprise IVR, chatbots, brand voice, healthcare	Excellent (Emotional Neural voices)	70+ languages	$15/M chars (Neural), Custom Voice extra	Cloud & On-prem, Full SDK
ElevenLabs	Commercial	Audiobooks, games, video narration, cloning	Exceptional (Human-like, expressive)	25,Äì30 languages	From $5/month, ~$733/M chars	Simple API, Cloud-only
Coqui TTS	Open Source	Offline apps, privacy-sensitive, custom languages	Good to Very Good (model-dependent)	1000+ languages (model-dependent)	Free (infra & dev cost only)	Self-hosted, full control

Agentic AI Tools

Agentic AI tools are like smart assistants that don't just wait around to be told what to do, they actually take initiative. You give them a goal, like "plan my trip" or "summarize these 10 reports and draft an email," and they figure out the steps, make decisions, and use different tools or data sources to get it done. Unlike regular AI that just responds to prompts, agentic AI can think through problems, adapt along the way, and even remember things from earlier conversations.

Rahul Kamath (https://www.linkedin.com/in/rahulnkamath/)

Tool Name	Type	Best Use Cases	Ease of Use	Integration Capability	Pricing Info	Model / LLM Used	Deployment	Autonomy Level
AutoGPT	Open Source	General-purpose autonomous task execution, research, content generation	Developer-friendly; requires setup and prompt tuning	Flexible, but dev-heavy (Python/Docker)	Free (MIT license); GPT-4 API costs extra (~$0.03 to $0.06 per 1K tokens)	GPT-4, GPT-3.5 (via OpenAI API)	Self-hosted	Fully autonomous
LangChain	Open Source	Custom agents with reasoning & memory, data-driven apps, RAG workflows	Developer-oriented; excellent docs & community	Extensive SDK + plugins; supports many APIs & data sources	Free core SDK (MIT); optional paid services (LangSmith, LangGraph)	Model-agnostic (supports GPT, Claude, Mistral, etc.)	Self-hosted / Cloud hybrid	Customizable (can do fully autonomous or hybrid)
IBM watsonx Orchestrate	Commercial	Enterprise workflow automation (HR, Sales, Procurement)	Business-user friendly (no-code flow builder)	Plug-and-play with enterprise tools (SAP, Workday, Outlook, etc.)	Starts ~$6K/month (Standard Edition); per-user pricing for scale	IBM Granite, OpenAI, Anthropic (pluggable via AI Gateway)	Cloud / On-prem	Semi to fully autonomous (pre-defined & user-configured agents)
Moveworks	Commercial	Employee support (IT, HR, Finance) via AI assistant	Extremely easy for employees; low-code for admins	Hundreds of built-in connectors + plugin system	Enterprise contract; based on users & volume (~6-7 figures/year)	Custom ensemble (can include GPT-4, Claude); private model layer	Cloud	Fully autonomous (with override); proactive actions

Rahul Kamath (https://www.linkedin.com/in/rahulnkamath/)